

# MTH4430 FinalProj Work

Mth 4330

Christopher Guzman, Dereck Lin, Haris Nioulikos, Samiha Uddin

2025-11-17

```
# Data Processing
student_health_data <- read.csv('Data/student_health_data.csv')
#convert response label to an ordered factor
student_health_data$Health_Risk_Level <- factor(student_health_data$Health_Risk_Level, levels = c("Low", "Moderate", "High"), ordered = TRUE)

#Encoding ordinal features, and the nominal feature gender
#cat features: Gender, Physical_Activity, Sleep_Quality, Mood
student_health_data <- student_health_data %>%
  mutate(
    Gender = case_when(
      Gender == "M" ~ 1,
      Gender == "F" ~ 0
    ),
    Physical_Activity = case_when(
      Physical_Activity == "High" ~ 2,
      Physical_Activity == "Moderate" ~ 1,
      Physical_Activity == "Low" ~ 0
    ),
    Sleep_Quality = case_when(
      Sleep_Quality == "Good" ~ 2,
      Sleep_Quality == "Moderate" ~ 1,
      Sleep_Quality == "Poor" ~ 0
    ),
    Mood = case_when(
      Mood == "Happy" ~ 2,
      Mood == "Neutral" ~ 1,
      Mood == "Stressed" ~ 0
    )
  )

# remove identifier(Student_ID)
student_health_data$Student_ID <- NULL

#OVR encode into separate cols
classes <- levels(student_health_data$Health_Risk_Level)
for(class in classes){
  isClass <- paste0("y_", class)
  student_health_data[[isClass]] <- ifelse(student_health_data$Health_Risk_Level == class, 1, 0)
}

#numerical encode for

set.seed(1)#for reproducibility

# split into train&test(8:2)
# src: https://scikit-learn.org/stable/common_pitfalls.html
idx <- createDataPartition(student_health_data$Health_Risk_Level, p = 0.8, list = FALSE)
tr <- student_health_data[idx,]
te <- student_health_data[-idx,]

#csv split assignments
```

```

v <- sample(1:5, nrow(tr), replace = TRUE)

#Baseline Model: Majority Rule by Christopher Guzman
ProjData <- student_health_data # To load the dataset
#print(df) # To view how the dataset would look like
dim(ProjData) # If curious about the dimensions of dataset

## [1] 1000   16

head(ProjData) # Want to see first few rows

##   Age Gender Heart_Rate Blood_Pressure_Systolic Blood_Pressure_Diastolic
## 1  24      1    50.66322            122.1730             84.41986
## 2  21      0    57.92604            110.7784             75.69615
## 3  22      1    59.29422            109.3757             83.80381
## 4  24      1    76.82623            125.1422             78.09159
## 5  20      1    68.34277            107.5156             80.67494
## 6  22      1    61.74415            90.0000             84.45086
##   Stress_Level_Biosensor Stress_Level_Self_Report Physical_Activity
## 1           3.137350          9.028669              2
## 2           3.699078          5.819697              1
## 3           6.785156          5.892360              0
## 4           6.408509          6.884001              2
## 5           7.264719          4.483450              1
## 6           4.262518          6.825001              1
##   Sleep_Quality Mood Study_Hours Project_Hours Health_Risk_Level y_Low
## 1            1     2    34.520973       16.80096 Moderate      0
## 2            2     0   16.763846       15.79115 Moderate      0
## 3            1     2    44.203798       25.67844 Moderate      0
## 4            0     2   21.776645       20.80839 High        0
## 5            0     2    8.964999       15.19405 Moderate      0
## 6            2     2    44.948229       15.65120 Moderate      0
##   y_Moderate y_High
## 1           1     0
## 2           1     0
## 3           1     0
## 4           0     1
## 5           1     0
## 6           1     0

#print(df$Health_Risk_Level) # To view the variable we are planning to predict
# names(ProjData) # To see the names of the columns of dataset

# str(ProjData) # if you want a summary of dataset
table(ProjData$Health_Risk_Level) # check the class distribution

## 
##   Low Moderate     High
##   190      672     138

# We can use the table() function to create frequency tables of categorical data of a certain vector.
# It will count the number of times a unique value or combination of values shows up in a vector.
# Can use for combination of vectors, creating a contingency table

ProjData$Health_Risk_Level = factor(ProjData$Health_Risk_Level) # Factor function used to handle categorical variables

# Baseline: Majority Class Predictor
# Find most frequent class
majority_class = names(which.max(table(ProjData$Health_Risk_Level)))
print(majority_class)

```

```

## [1] "Moderate"

# We will predict this class for every student
baseline_predictions = rep(majority_class, nrow(ProjData))

# Compute Accuracy
baseline_accuracy = mean(baseline_predictions == ProjData$Health_Risk_Level)
print(baseline_accuracy)

```

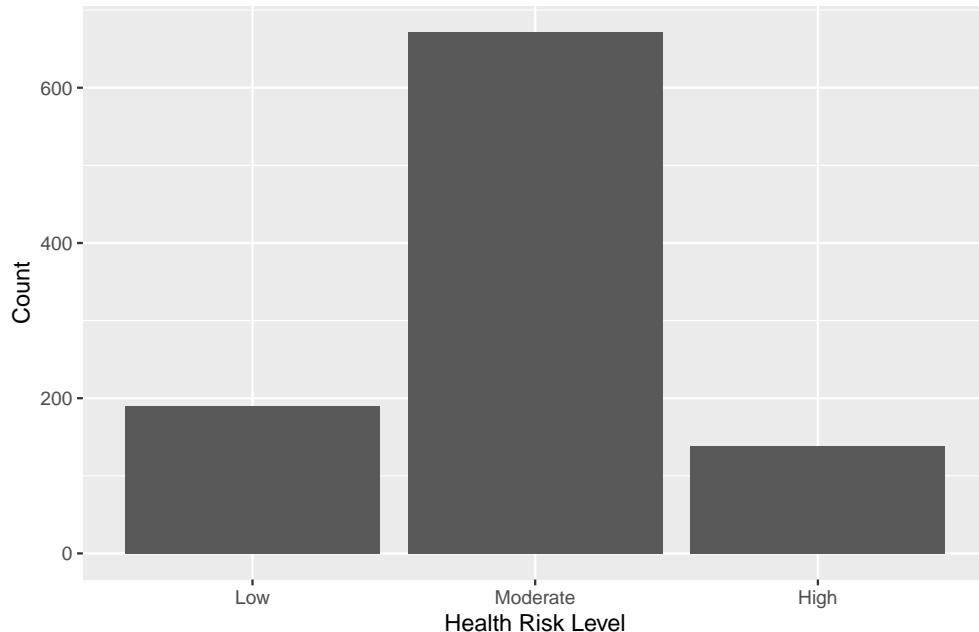
```
## [1] 0.672
```

```

ggplot(ProjData, aes(x = Health_Risk_Level)) +
  geom_bar() +
  labs(title = "Class Distribution of Health Risk Level",
       x = "Health Risk Level",
       y = "Count")

```

Class Distribution of Health Risk Level



## Base ovr Log Reg Model

```

# predictors: "Age + Gender + Heart_Rate + Blood_Pressure_Systolic + Blood_Pressure_Diastolic + Stress_Level_Biosensor + Stress_L
logr_cv_log_loss <- c(0,0,0,0,0)

#itr over cv folds
for(fold in 1:5){
  #split tr data into folds
  cv_tr <- tr[v != fold, ]
  cv_te <- tr[v == fold, ]

  #fit on all, excluding response labels(Health_Risk_Level, y_Low, y_Moderate, y_High)
  logr_High <- glm(y_High ~ Age + Gender + Heart_Rate + Blood_Pressure_Systolic + Blood_Pressure_Diastolic + Stress_Level_Biosen
  logr_Moderate <- glm(y_Moderate ~ Age + Gender + Heart_Rate + Blood_Pressure_Systolic + Blood_Pressure_Diastolic + Stress_Leve
  logr_Low <- glm(y_Low ~ Age + Gender + Heart_Rate + Blood_Pressure_Systolic + Blood_Pressure_Diastolic + Stress_Level_Biosensc

  # #predict

```

```

logr_High_Pred <- predict(logr_High, newdata = cv_te, type = 'response')
logr_Moderate_Pred <- predict(logr_Moderate, newdata = cv_te, type = 'response')
logr_Low_Pred <- predict(logr_Low, newdata = cv_te, type = 'response')

#normalize probabilities
logr_y_pred <- cbind(logr_Low_Pred,logr_Moderate_Pred,logr_High_Pred)
logr_y_pred_row_sum <- rowSums(logr_y_pred)
logr_y_pred_normalized <- logr_y_pred / logr_y_pred_row_sum

#real label
logr_y_real <- cv_te[,c("y_Low","y_Moderate","y_High")]

#Cross entropy loss
logr_cv_log_loss[fold] <- -mean(rowSums(logr_y_real * log(logr_y_pred_normalized)))
}

logr_log_loss <- mean(logr_cv_log_loss)
logr_log_loss

## [1] 0.4337477

```

## Refitting base ovr Logr on entirety of train set

```

final_logr_High <- glm(y_High ~ Age + Gender + Heart_Rate + Blood_Pressure_Systolic + Blood_Pressure_Diastolic + Stress_Level_Bios
                        data = tr, family = "binomial")

final_logr_Moderate <- glm(y_Moderate ~ Age + Gender + Heart_Rate + Blood_Pressure_Systolic + Blood_Pressure_Diastolic + Stress_Level_Bios
                            data = tr, family = "binomial")

final_logr_Low <- glm(y_Low ~ Age + Gender + Heart_Rate + Blood_Pressure_Systolic + Blood_Pressure_Diastolic + Stress_Level_Bios
                        data = tr, family = "binomial")

```

## Pred Final ovr Log R model on Train Set

```
## [1] 0.4067402
```

## Pred Final ovr Log R model on Test Set

```

#Predict
final_logr_High_pred <- predict(final_logr_High, newdata = te, type = 'response')
final_logr_Moderate_pred <- predict(final_logr_Moderate, newdata = te, type = 'response')
final_logr_Low_pred <- predict(final_logr_Low, newdata = te, type = 'response')

# bind into prob vec
final_logr_pred_prob <- cbind(final_logr_Low_pred,final_logr_Moderate_pred,final_logr_High_pred)
final_logr_pred_prob_rsum <- rowSums(final_logr_pred_prob)
final_logr_pred_prob_normalized <- final_logr_pred_prob / final_logr_pred_prob_rsum

final_logr_log_loss <- -mean(rowSums(te[,c("y_Low","y_Moderate","y_High")] * log(final_logr_pred_prob_normalized)))
final_logr_log_loss

## [1] 0.4616341

```

## Sanity check, comparing Majority rule with Logr

```

#Vector of hard classes extracted from test set
y_te_real <- te[, "Health_Risk_Level"]

#collapsing 'logr_y_pred_normalized' into hard classes
logr_y_pred_class <- max.col(final_logr_pred_prob_normalized)

#using accuracy as a baseline metric
desired_levels <- c("Low", "Moderate", "High")
colnames(final_logr_pred_prob_normalized) <- desired_levels
final_predictions <- colnames(final_logr_pred_prob_normalized)[logr_y_pred_class]

final_predictions <- factor(final_predictions, levels = desired_levels)

logr_conf_matrix <- table(Predicted = final_predictions, Actual = te$Health_Risk_Level)

confusionMatrix(logr_conf_matrix, mode = "everything")

## Confusion Matrix and Statistics
##
##           Actual
## Predicted   Low Moderate High
##   Low        31      1     3
##   Moderate    7     130    14
##   High        0      3    10
##
## Overall Statistics
##
##           Accuracy : 0.8593
##             95% CI : (0.8031, 0.9044)
##   No Information Rate : 0.6734
##   P-Value [Acc > NIR] : 1.644e-09
##
##           Kappa : 0.6849
##
## McNemar's Test P-Value : 0.002174
##
## Statistics by Class:
##
##           Class: Low Class: Moderate Class: High
## Sensitivity          0.8158          0.9701          0.37037
## Specificity           0.9752          0.6769          0.98256
## Pos Pred Value        0.8857          0.8609          0.76923
## Neg Pred Value        0.9573          0.9167          0.90860
## Precision            0.8857          0.8609          0.76923
## Recall               0.8158          0.9701          0.37037
## F1                   0.8493          0.9123          0.50000
## Prevalence            0.1910          0.6734          0.13568
## Detection Rate        0.1558          0.6533          0.05025
## Detection Prevalence  0.1759          0.7588          0.06533
## Balanced Accuracy     0.8955          0.8235          0.67646

```

## ovr LASSO

```

lasso_cv_log_loss <- c(0,0,0,0,0)
for(fold in 1:5){
  #13-16 = health risk lvl, y_...,y_...,y_...
  #split tr data into folds and extract feature matrix and response matrix for lasso using glmnet, also excluding hard class labels
  cv_x_tr <- as.matrix(tr[v != fold, - c(13,14,15,16)])
  cv_y_tr <- tr[v != fold, c("y_Low", "y_Moderate", "y_High")]
  cv_x_te <- as.matrix(tr[v == fold, - c(13,14,15,16)])
  cv_y_te <- tr[v == fold,c("y_Low", "y_Moderate", "y_High")]

  #train

```

```

lasso_high <- cv.glmnet(cv_x_tr, cv_y_tr[, "y_High"], family = "binomial", type.measure = "deviance")
lasso_moderate <- cv.glmnet(cv_x_tr, cv_y_tr[, "y_Moderate"], family = "binomial", type.measure = "deviance")
lasso_low <- cv.glmnet(cv_x_tr, cv_y_tr[, "y_Low"], family = "binomial", type.measure = "deviance")

#pred
lasso_high_pred <- predict(lasso_high, newx = cv_x_te, s = "lambda.min", type = "response")
lasso_moderate_pred <- predict(lasso_moderate, newx = cv_x_te, s = "lambda.min", type = "response")
lasso_low_pred <- predict(lasso_low, newx = cv_x_te, s = "lambda.min", type = "response")

#prob matrix
lasso_y_pred <- cbind(lasso_low_pred, lasso_moderate_pred, lasso_high_pred)
lasso_y_pred_rowsums <- rowSums(lasso_y_pred)
lasso_y_pred_normalized <- lasso_y_pred / lasso_y_pred_rowsums

#Log loss calc
lasso_cv_log_loss[fold] <- -mean(rowSums(cv_y_te * log(lasso_y_pred_normalized)))
}

lasso_log_loss <- mean(lasso_cv_log_loss)
lasso_log_loss

## [1] 0.4404227

```

## LASSO refit on Train set

```

X_tr_full_matrix <- as.matrix(tr[, -c(13, 14, 15, 16)])
Y_tr_full_targets <- tr[, c(14, 15, 16)]

best_lasso_high <- cv.glmnet(X_tr_full_matrix, tr[, 16], family = "binomial", type.measure = "deviance")
best_lasso_moderate <- cv.glmnet(X_tr_full_matrix, tr[, 15], family = "binomial", type.measure = "deviance")
best_lasso_low <- cv.glmnet(X_tr_full_matrix, tr[, 14], family = "binomial", type.measure = "deviance")

best_lasso_high_pred <- predict(best_lasso_high, newx = X_tr_full_matrix, s = "lambda.min", type = "response")
best_lasso_moderate_pred <- predict(best_lasso_moderate, newx = X_tr_full_matrix, s = "lambda.min", type = "response")
best_lasso_low_pred <- predict(best_lasso_low, newx = X_tr_full_matrix, s = "lambda.min", type = "response")

best_lasos_y_pred <- cbind(best_lasso_low_pred, best_lasso_moderate_pred, best_lasso_high_pred)
best_lasso_y_pred_rowsum <- rowSums(best_lasos_y_pred)
best_lasso_y_pred_normalized <- best_lasos_y_pred / best_lasso_y_pred_rowsum

log_loss <- -mean(rowSums(Y_tr_full_targets * log(best_lasso_y_pred_normalized)))
log_loss

## [1] 0.4244026

best_lambda <- c(best_lasso_low$lambda.min, best_lasso_moderate$lambda.min, best_lasso_high$lambda.min)
best_lambda

## [1] 0.004015451 0.009399397 0.006117351

```

## Optimized Lasso Model

```

X_te_full_matrix <- as.matrix(te[, -c(13, 14, 15, 16)])
Y_te_full_targets <- te[, c(14, 15, 16)]

te_best_lasso_high_pred <- predict(best_lasso_high, newx = X_te_full_matrix, s = best_lambda[3], type = "response")
te_best_lasso_moderate_pred <- predict(best_lasso_moderate, X_te_full_matrix, s = best_lambda[2], type = "response")
te_best_lasso_low_pred <- predict(best_lasso_low, newx = X_te_full_matrix, s = best_lambda[1], type = "response")

```

```

te_best_lasso_y_pred <- cbind(te_best_lasso_low_pred, te_best_lasso_moderate_pred, te_best_lasso_high_pred)
te_best_lasso_y_pred_rsum <- rowSums(te_best_lasso_y_pred)
te_best_lasso_y_pred_normalized <- te_best_lasso_y_pred / te_best_lasso_y_pred_rsum

te_best_lasso_log_loss <- -mean(rowSums(Y_te_full_targets * log(te_best_lasso_y_pred_normalized)))
te_best_lasso_log_loss

## [1] 0.4676428

## Optimized Lasso conf mat

#y_te_real = true hard classes
#collapsing 'te_best_lasso_y_pred_normalized' into hard classes
colnames(te_best_lasso_y_pred_normalized) <- c("Low", "Moderate", "High")
best_lasso_y_pred_class <- max.col(te_best_lasso_y_pred_normalized)

final_predictions <- colnames(te_best_lasso_y_pred_normalized)[best_lasso_y_pred_class]

final_predictions <- factor(final_predictions, levels = desired_levels)

lasso_conf_matrix <- table(Predicted = final_predictions, Actual = te$Health_Risk_Level)

confusionMatrix(lasso_conf_matrix, mode = "everything")

## Confusion Matrix and Statistics
##
##          Actual
## Predicted  Low Moderate High
##   Low      28       1     3
##   Moderate 10     132    18
##   High      0       1     6
##
## Overall Statistics
##
##          Accuracy : 0.8342
##             95% CI : (0.7751, 0.883)
##    No Information Rate : 0.6734
##    P-Value [Acc > NIR] : 2.382e-07
##
##          Kappa : 0.6081
##
## McNemar's Test P-Value : 1.171e-05
##
## Statistics by Class:
##
##          Class: Low Class: Moderate Class: High
## Sensitivity          0.7368      0.9851      0.22222
## Specificity           0.9752      0.5692      0.99419
## Pos Pred Value        0.8750      0.8250      0.85714
## Neg Pred Value        0.9401      0.9487      0.89062
## Precision            0.8750      0.8250      0.85714
## Recall               0.7368      0.9851      0.22222
## F1                   0.8000      0.8980      0.35294
## Prevalence            0.1910      0.6734      0.13568
## Detection Rate        0.1407      0.6633      0.03015
## Detection Prevalence  0.1608      0.8040      0.03518
## Balanced Accuracy     0.8560      0.7772      0.60820

confusionMatrix(logr_conf_matrix, mode = "everything")

## Confusion Matrix and Statistics
##
##          Actual

```

```

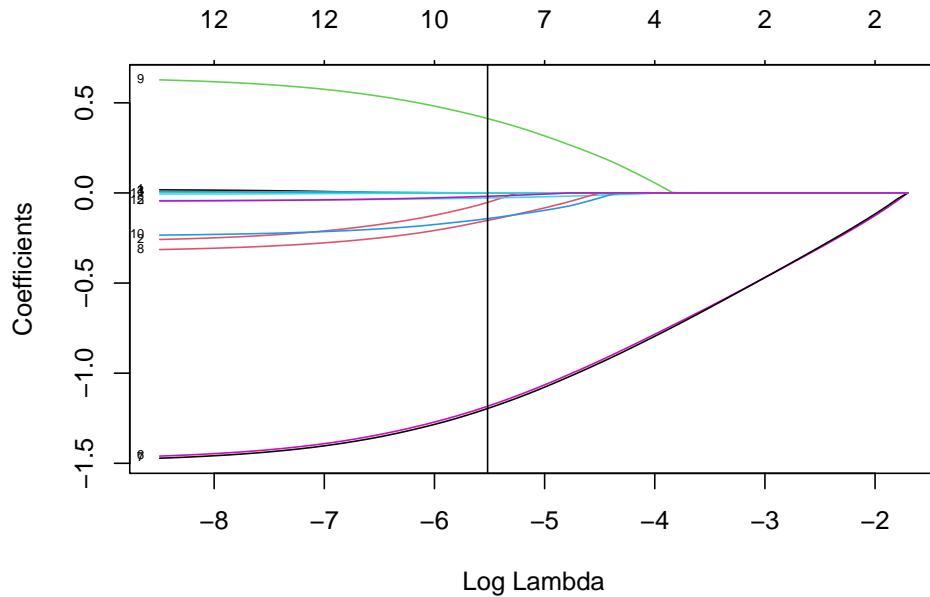
## Predicted Low Moderate High
##   Low      31       1     3
## Moderate    7     130    14
## High       0       3    10
##
## Overall Statistics
##
##           Accuracy : 0.8593
##         95% CI : (0.8031, 0.9044)
##    No Information Rate : 0.6734
## P-Value [Acc > NIR] : 1.644e-09
##
##           Kappa : 0.6849
##
## McNemar's Test P-Value : 0.002174
##
## Statistics by Class:
##
##           Class: Low Class: Moderate Class: High
## Sensitivity          0.8158          0.9701          0.37037
## Specificity          0.9752          0.6769          0.98256
## Pos Pred Value       0.8857          0.8609          0.76923
## Neg Pred Value       0.9573          0.9167          0.90860
## Precision            0.8857          0.8609          0.76923
## Recall               0.8158          0.9701          0.37037
## F1                   0.8493          0.9123          0.50000
## Prevalence            0.1910          0.6734          0.13568
## Detection Rate        0.1558          0.6533          0.05025
## Detection Prevalence 0.1759          0.7588          0.06533
## Balanced Accuracy     0.8955          0.8235          0.67646

```

```

plot(best_lasso_low$glmnet.fit, xvar = "lambda", label = TRUE)
abline(v = log(best_lambda)[1])

```



```

coef(best_lasso_low, s = "lambda.min")

```

```

## 13 x 1 sparse Matrix of class "dgCMatrix"

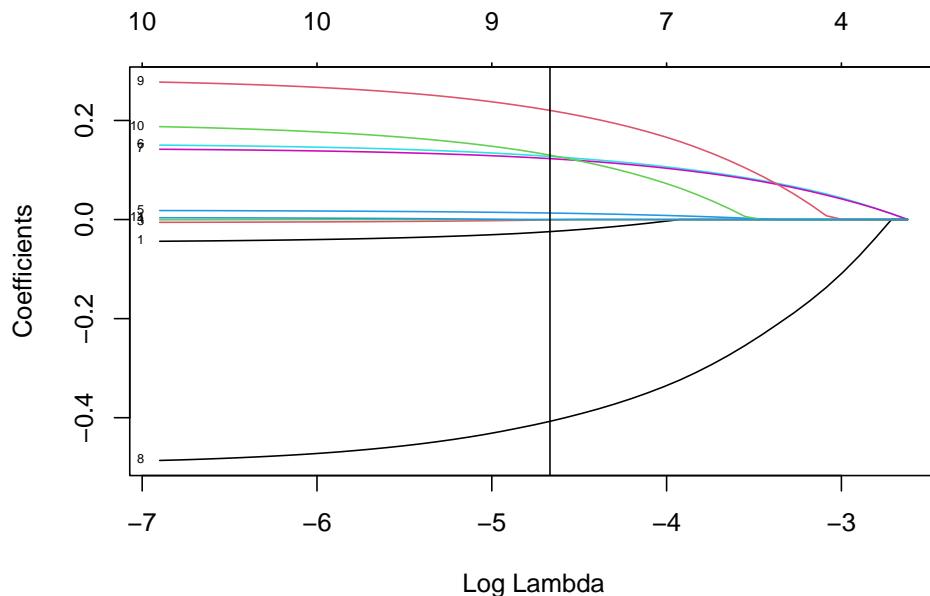
```

```

##                               s1
## (Intercept)      10.91029653
## Age              .
## Gender           -0.05234152
## Heart_Rate       .
## Blood_Pressure_Systolic .
## Blood_Pressure_Diastolic -0.02768900
## Stress_Level_Biosensor -1.18340279
## Stress_Level_Self_Report -1.19615320
## Physical_Activity -0.15274027
## Sleep_Quality    0.41269211
## Mood              -0.14247052
## Study_Hours      .
## Project_Hours   -0.01861177

plot(best_lasso_moderate$glmnet.fit, xvar = "lambda", label = TRUE)
abline(v = log(best_lambda)[2])

```



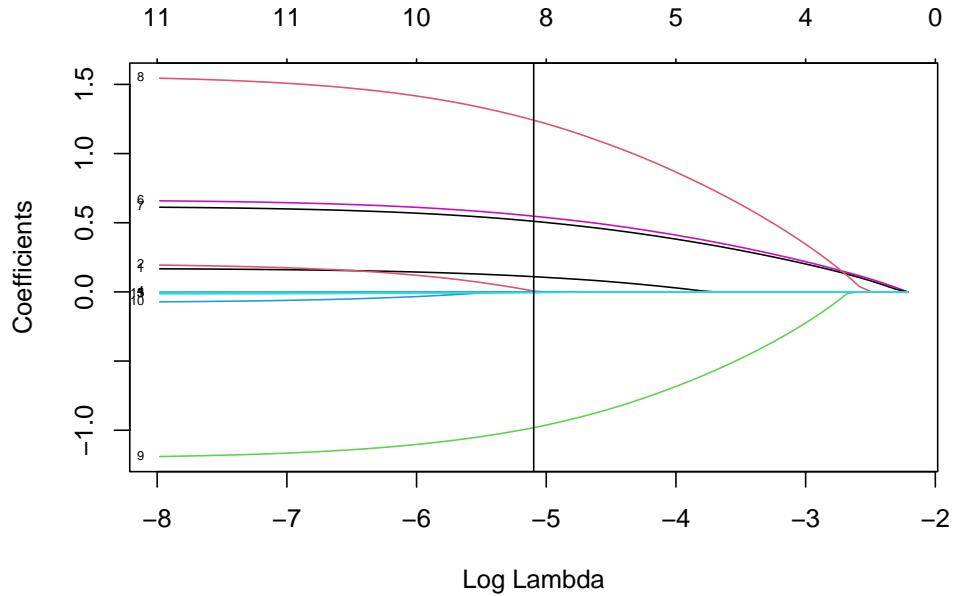
```
coef(best_lasso_moderate, s = "lambda.min")
```

```

## 13 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept)      -1.1382434826
## Age             -0.0243767507
## Gender           .
## Heart_Rate      -0.0007822639
## Blood_Pressure_Systolic .
## Blood_Pressure_Diastolic 0.0131325741
## Stress_Level_Biosensor 0.1273952093
## Stress_Level_Self_Report 0.1229408006
## Physical_Activity -0.4073551144
## Sleep_Quality    0.2203588768
## Mood             0.1299367738
## Study_Hours      .
## Project_Hours   .

```

```
plot(best_lasso_high$glmnet.fit, xvar = "lambda", label = TRUE)
abline(v = log(best_lambda)[3])
```



```
coef(best_lasso_high, s = "lambda.min")
```

```
## 13 x 1 sparse Matrix of class "dgCMatrix"
##           s1
## (Intercept) -10.441838638
## Age          0.109930955
## Gender        0.005887929
## Heart_Rate   -0.001298608
## Blood_Pressure_Systolic .
## Blood_Pressure_Diastolic -0.004741060
## Stress_Level_Biosensor  0.547092707
## Stress_Level_Self_Report 0.510406793
## Physical_Activity      1.241046947
## Sleep_Quality         -0.981167905
## Mood                 .
## Study_Hours          -0.004480365
## Project_Hours        .
```